Introducción a la Bioinformática:

Hidden Markov Models (HMMs)

Luis Garreta

Doctorado en Ingeniería
Pontificia Universidad Javeriana – Cali

March 29, 2017

Applications of HMMs

Gene Finding and Prediction

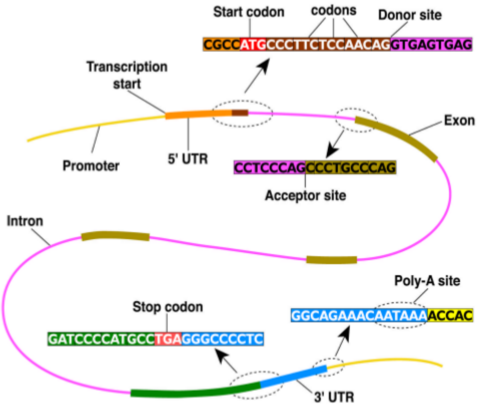

Protein- Profile HMMs

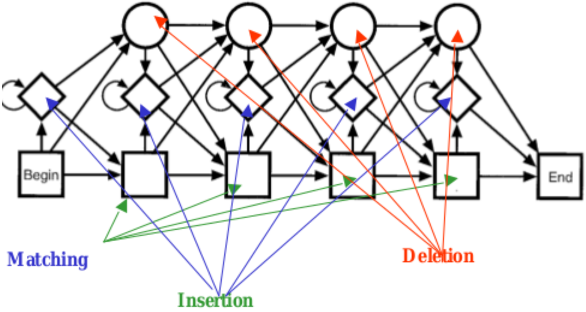

Protein- Profile HMMs

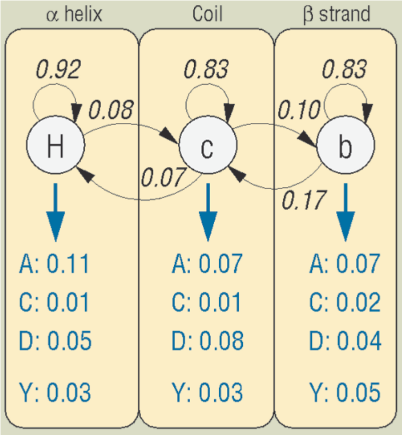

Other Applications of HMMs

- Speech recognition
- Optical character recognition
- Spell checking

Hidden Markov Model (HMM) Architecture

Markov Chains:

Markov Assumption

Three states of weather

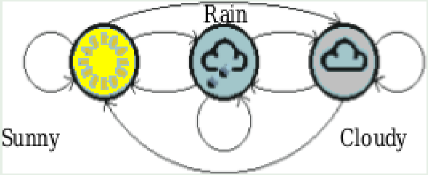

- Three states: Sunny, Cloudy, and Rainy
- Weather pattern instead deterministically

- Markov Assumption: The state of the model depends only upon the previous states of the model
- Order n Model (First Order): The choice of state is made purely on the basis of the previous state

Markov Chains:

State Transition Matrix (A)

$$A = \begin{array}{cc} & \text{Today} \\ \text{Yesterday} & \begin{array}{c|ccc} & \text{Sun} & \text{Cloud} & \text{Rain} \\ \hline \text{Sun} & 0.5 & 0.25 & 0.25 \\ \text{Cloud} & 0.375 & 0.125 & 0.375 \\ \text{Rain} & 0.125 & 0.625 & 0.375 \end{array} \end{array}$$

It it was sunny yesterday, there is a probability of 0.5 that it will be sunny today, and 0.25 that it will be cloudy or rainy.

Markov Chains:

Vector of Initial Probabilities (π)

$$\Pi = \begin{array}{ccc} \text{Sun} & \text{Cloud} & \text{Rain} \\ \end{array}$$
$$\begin{pmatrix} 1.0 & 0.0 & 0.0 \end{pmatrix}$$

- To initialize such a system, we need to state what the weather was (or probably was) on the day after creation;
- So, we know it was sunny on day 1

Markov Chains:

First Order Markov Process

- States: Three states: sunny, cloudy, rainy
- π vector: Probability of the system in each states at time 0
- State transition Matrix: Probability of the weather given the previous day's weather

Any system that can be described in this manner is a Markov process.

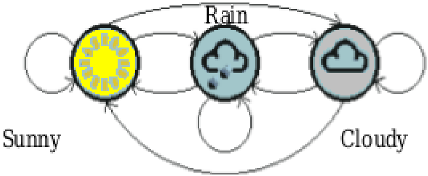

$$A = \begin{array}{c} \\ \text{Sun} \\ \text{Cloud} \\ \text{Rain} \end{array} \begin{array}{c} \begin{array}{ccc} \text{Sun} & \text{Cloud} & \text{Rain} \end{array} \\ \begin{pmatrix} 0.5 & 0.25 & 0.25 \\ 0.375 & 0.125 & 0.375 \\ 0.125 & 0.625 & 0.375 \end{pmatrix} \end{array}$$

$$\Pi = \begin{array}{c} \begin{array}{ccc} \text{Sun} & \text{Cloud} & \text{Rain} \end{array} \\ \begin{pmatrix} 1.0 & 0.0 & 0.0 \end{pmatrix} \end{array}$$

Hidden Markov Models

- In some cases the patterns that we wish to find are not described sufficiently by a Markov process.
- A hermit for instance may not have access to direct weather observations, but does have a piece of seaweed.

- Sea and weather lore: seaweeds are weather predictors (they absorb atmospheric moisture)
- The seaweed is probabilistically related to the state of the weather:

Hidden Markov Models: Two sets of States

In this case we have two sets of states:

- observable states (the state of the seaweed) and
- hidden states (the state of the weather).

We wish to devise an algorithm for the hermit to forecast weather from the seaweed and the Markov assumption without actually ever seeing the weather.

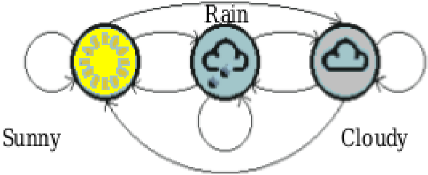

Hidden Markov Models: Hidden and Observable States

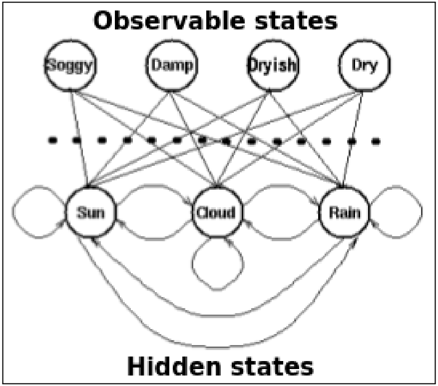

- Hidden states (the true weather) are modeled by a simple Markov process.
- So, they are all connected to each other.
- The new connections represent: *the probability of generating a particular observed state given that the Markov process is in a particular hidden state.*

Hidden Markov Models: Emission Matrix

The probabilities of the observable states given a particular hidden state:

Seaweed

Weather	Dry	Dryish	Damp	Soggy
Sun	0.60	0.20	0.15	0.05
Cloud	0.25	0.25	0.25	0.25
Rain	0.05	0.10	0.35	0.50

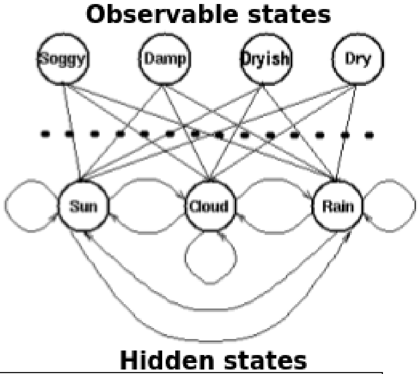

All probabilities "entering" an observable state will sum to 1 :

$$Pr(Obs|Sun) + Pr(Obs|Cloud) + Pr(Obs|Rain) = 1$$

Example: The Dishonest Casino

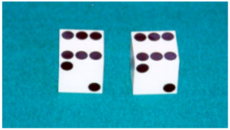

Game:

1. You bet $1
2. You roll
3. Casino player rolls
4. Highest number wins $2

The casino has two dice:

Fair die

P(1) = P(2) = P(3) = P(5) = P(6) = 1/6

Loaded die

P(1) = P(2) = P(3) = P(5) = 1/10

P(6) = 1/2

Casino player switches between fair and loaded die (not too often, and not for too long)

The dishonest casino model

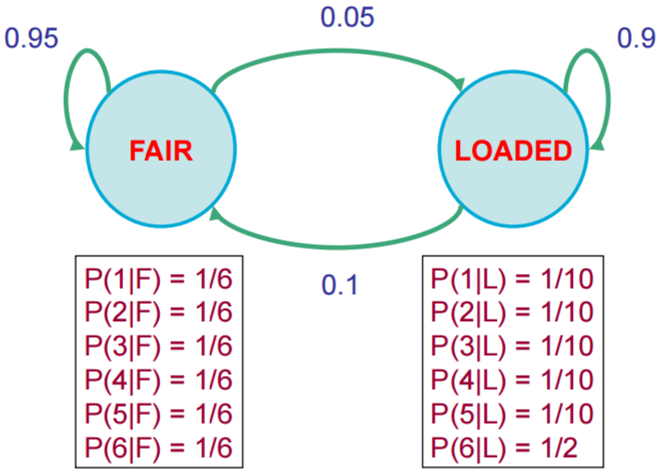

Question # 1 – Evaluation

GIVEN:

A sequence of rolls by the casino player

$$\underbrace{1245526462146146136136661664661636616366163616515615115146123562344}_{\boxed{\text{Prob} = 1.3 \times 10^{-35}}}$$

QUESTION:

How likely is this sequence, given our model of how the casino works?

This is the **EVALUATION** problem in HMMs

Question # 2 – Decoding

GIVEN:

A sequence of rolls by the casino player

124552646214614613	613666166466163661636616361616	515615115146123562344
FAIR	LOADED	FAIR

QUESTION:

What portion of the sequence was generated with the fair die, and what portion with the loaded die?

This is the **DECODING** question in HMMs

Question # 3 – Learning

GIVEN:

A sequence of rolls by the casino player

1245526462146146136136661664661636616366163616515615115146123562344

QUESTION:

How does the casino player work: How "loaded" is the loaded die? How "fair" is the fair die? How often does the casino player change from fair to loaded, and back?

This is the **LEARNING** question in HMMs

The dishonest casino model

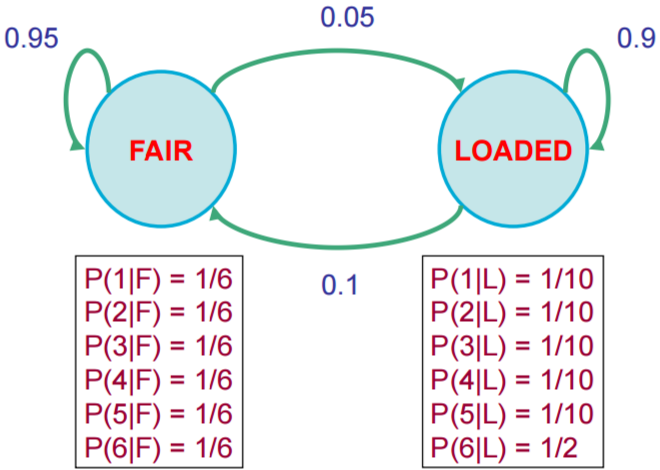

Definition of a hidden Markov model

- Alphabet $\Sigma = \{ b_1, b_2, \ldots, b_M \}$
- Set of states $Q = \{ 1, \ldots, K \}$ $(K = |Q|)$
- Transition probabilities between any two states

  a_{ij} = transition probability from state i to state j

  $a_{i1} + \ldots + a_{iK} = 1$, for all states i
- Initial probabilities a_{0i}

  $a_{01} + \ldots + a_{0K} = 1$
- Emission probabilities within each state

  $e_k(b) = P(x_i = b \mid \pi_i = k)$

  $e_k(b_1) + \ldots + e_k(b_M) = 1$

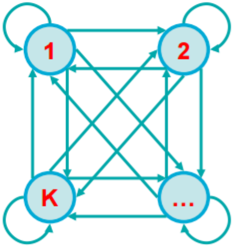

Hidden states and observed sequence

At time step t,

π_t denotes the (hidden) state in the Markov chain

x_t denotes the symbol emitted in state π_t

A path of length N is: $\pi_1, \pi_2, \ldots, \pi_N$

An observed sequence of length N is: $x_1, x_2, \ldots, x_N$

An HMM is "memory-less"

At time step t, the only thing that affects the next state is the current State, π_t

$P(\pi_{t+1} = k \mid$ "whatever happened so far")
$= P(\pi_{t+1} = k \mid \pi_1, \pi_2, \ldots, \pi_t, x_1, x_2, \ldots, x_t)$
$= P(\pi_{t+1} = k \mid \pi_t)$

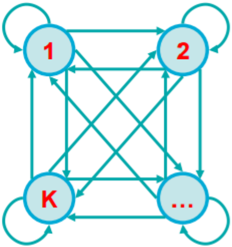

A parse of a sequence

Given a sequence $x = x_1 \ldots\ldots x_N$,

A parse of x is a sequence of states $\pi = \pi_1, \ldots\ldots, \pi_N$

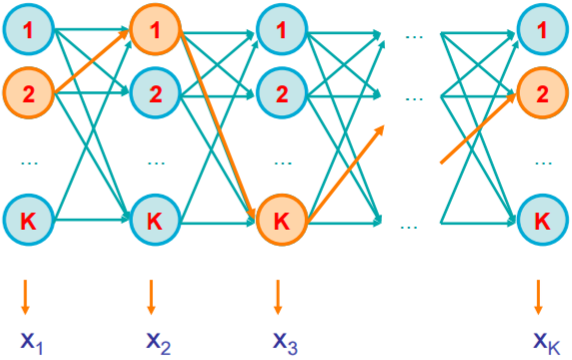

Likelihood of a parse: $P(x, \pi)$

Given a sequence $x = x_1 \ldots\ldots, x_N$
and a parse $\pi = \pi_1, \ldots\ldots, \pi_N$,

How likely is the parse
(given our HMM)?

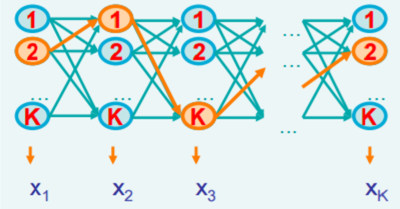

$$
\begin{aligned}
P(x, \pi) &= P(x_1, \ldots, x_N, \pi_1, \ldots\ldots, \pi_N) \\
&= P(x_N, \pi_N \mid x_1 \ldots x_{N-1}, \pi_1, \ldots\ldots, \pi_{N-1})\, P(x_1 \ldots x_{N-1}, \pi_1, \ldots\ldots, \pi_{N-1}) \\
&= P(x_N, \pi_N \mid \pi_{N-1})\, P(x_1 \ldots x_{N-1}, \pi_1, \ldots\ldots, \pi_{N-1}) \\
&= \ldots \\
&= P(x_N, \pi_N \mid \pi_{N-1})\, P(x_{N-1}, \pi_{N-1} \mid \pi_{N-2}) \ldots\ldots P(x_2, \pi_2 \mid \pi_1)\, P(x_1, \pi_1) \\
&= P(x_N \mid \pi_N)\, P(\pi_N \mid \pi_{N-1}) \ldots\ldots P(x_2 \mid \pi_2)\, P(\pi_2 \mid \pi_1)\, P(x_1 \mid \pi_1)\, P(\pi_1) \\
&= a_{0\pi 1}\, a_{\pi 1 \pi 2} \ldots\ldots a_{\pi N-1 \pi N}\, e_{\pi 1}(x_1) \ldots\ldots e_{\pi N}(x_N) \\
&= \prod_{i=1}^{N} a_{\pi_{i-1}\pi_i} e_{\pi_i}(x_i)
\end{aligned}
$$

Example: the dishonest casino $P(x, \pi)$, $\pi = FFFFFF...FF$

What is the probability of a sequence of rolls

x = 1, 2, 1, 5, 6, 2, 1, 6, 2, 4

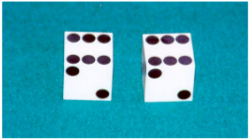

and the parse

π = Fair, Fair, Fair, Fair, Fair, Fair, Fair, Fair, Fair, Fair?

(say initial probs $a_{0,Fair} = ½$, $a_{0,Loaded} = ½$)

½ × P(1 | Fair) P(Fair | Fair) P(2 | Fair) P(Fair | Fair) … P(4 | Fair) =

$½ \times (1/6)^{10} \times (0.95)^9 = 5.2 \times 10^{-9}$

Example: the dishonest casino $P(x, \pi)$, $\pi = LLLLLL...LL$

So, the likelihood the die is fair in all this run is 5.2×10^{-9}

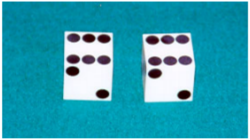

What about

π = Loaded, Loaded, Loaded, Loaded, Loaded, Loaded, Loaded, Loaded, Loaded, Loaded?

½ × P(1 | Loaded) P(Loaded|Loaded) … P(4 | Loaded) =

$\frac{1}{2} \times (1/10)^8 \times (1/2)^2 (0.9)^9 = 4.8 \times 10^{-10}$

Therefore, it is more likely that the die is fair all the way, than loaded all the way

Example: the dishonest casino, loglikehood-ratio

A likelihood ratio test is a statistical test used for comparing the goodness of fit of two models, one of which (the null model) is a special case of the other (the alternative model)

$$log(\frac{P(X|\pi_{Fair})}{P(X|\pi_{Loaded})}) = log(\frac{5.2-09}{4.8e-10}) = 10.76$$

Example: the dishonest casino: Suspicion of loaded dice

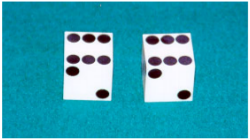

Let the sequence of rolls be:

x = 1, 6, 6, 5, 6, 2, 6, 6, 3, 6

And let's consider π = F, F,..., F

$P(x, \pi) = ½ \times (1/6)^{10} \times (0.95)^9 = 5.2 \times 10^{-9}$
(same as before)

And for π = L, L,..., L:

$P(x, \pi) = ½ \times (1/10)^4 \times (1/2)^6 (0.9)^9 = 3.02 \times 10^{-7}$

So, the observed sequence is ~100 times more likely if a loaded die is used

Clarification of notation

P[x | M]: The probability that sequence x was generated by the model

The model is: architecture (#states, etc)
+ parameters $\theta = a_{ij}, e_i(.)$

So, P[x | M] is the same as P[x | θ], and P[x], when the architecture, and the parameters, respectively, are implied

Similarly, P[x, π | M], P[x, π | θ] and P[x, π] are the same when the architecture, and the parameters, are implied

In the LEARNING problem we write P[x | θ] to emphasize that we are seeking the θ* that maximizes P[x | θ]

What we know

Given a sequence $x = x_1, \dots, x_N$
and a parse $\pi = \pi_1, \dots, \pi_N,$

we know how to compute
how likely the parse is:

$P(x, \pi)$

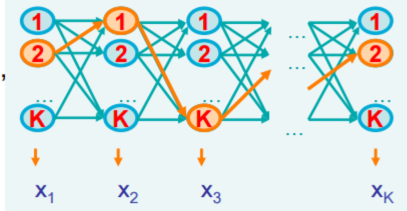

What we would know

1. Evaluation

   GIVEN HMM M, and a sequence x,
   FIND Prob[x | M]

2. Decoding

   GIVEN HMM M, and a sequence x,
   FIND the sequence π of states that maximizes P[x, π | M]

3. Learning

   GIVEN HMM M, with unspecified transition/emission probs., and a sequence x,
   FIND parameters $\theta = (e_i(.), a_{ij})$ that maximize P[x | θ]